

Final Year B. Tech., Sem VII 2022-23

High Performance Computing Lab

PRN: 2020BTECS00206

Full Name: SAYALI YOGESH DESAI

Batch: B4

Assignment No. 9

-
- 1. Implement Vector-Vector addition using CUDA C. State and justify the speedup using different size of threads and blocks.**

```
#include <stdio.h>

/*
 * Host function to initialize vector elements. This function
 * simply initializes each element to equal its index in the
 * vector.
 */

void initWith(float num, float *a, int N)
{
    for(int i = 0; i < N; ++i)
    {
        a[i] = num;
    }
}

/*
 * Device kernel stores into `result` the sum of each
 * same-indexed value of `a` and `b`.
 */
```

```

__global__
void addVectorsInto(float *result, float *a, float *b, int N)
{
    int index = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x * gridDim.x;

    for(int i = index; i < N; i += stride)
    {
        result[i] = a[i] + b[i];
    }
}

/*
 * Host function to confirm values in `vector`. This function
 * assumes all values are the same `target` value.
 */

void checkElementsAre(float target, float *vector, int N)
{
    for(int i = 0; i < N; i++)
    {
        if(vector[i] != target)
        {
            printf("FAIL: vector[%d] - %0.0f does not equal %0.0f\n", i, vector[i], target);
            exit(1);
        }
    }
    printf("Success! All values calculated correctly.\n");
}

int main()
{

```

```

const int N = 2<<24;
size_t size = N * sizeof(float);

float *a;
float *b;
float *c;

cudaMallocManaged(&a, size);
cudaMallocManaged(&b, size);
cudaMallocManaged(&c, size);

initWith(3, a, N);
initWith(4, b, N);
initWith(0, c, N);

size_t threadsPerBlock;
size_t numberOfBlocks;

/*
 * nsys should register performance changes when execution configuration
 * is updated.
 */

threadsPerBlock = 1;
numberOfBlocks = 1;

cudaError_t addVectorsErr;
cudaError_t asyncErr;

addVectorsInto<<<numberOfBlocks, threadsPerBlock>>>(c, a, b, N);

addVectorsErr = cudaGetLastError();

```

```

if(addVectorsErr != cudaSuccess) printf("Error: %s\n",
cudaGetErrorString(addVectorsErr));

```

```

asyncErr = cudaDeviceSynchronize();
if(asyncErr != cudaSuccess) printf("Error: %s\n", cudaGetErrorString(asyncErr));

```

```

checkElementsAre(7, c, N);

```

```

cudaFree(a);

```

```

cudaFree(b);

```

```

cudaFree(c);

```

```

}

```

N=2<<16

Serial Execution Time: 17944640

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
93.2	254696095	3	84898698.3	4141	254670295	cudaMallocManaged
6.6	17955956	1	17955956.0	17955956	17955956	cudaDeviceSynchronize
0.2	566454	3	188818.0	33156	416960	cudaFree
0.0	39346	1	39346.0	39346	39346	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	17944640	1	17944640.0	17944640	17944640	addVectorsInto(float*, float*, float*, int)

Parallel Execution Time:

Number of Blocks	Threads per Block	Time (in ns)	Speedup
8	256	1305172	13.7488
16	256	1309816	13.7001
32	256	1406457	12.7587
8	512	1567085	11.4509
16	512	1391095	12.8996
132	512	1384345	12.9625

Number of blocks: 8, Threads per block: 256, Time Required: 1305172

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.3	242232114	3	80744038.0	3776	242214856	cudaMallocManaged
0.5	1308578	1	1308578.0	1308578	1308578	cudaDeviceSynchronize
0.1	254964	3	84988.0	10490	212916	cudaFree
0.0	35428	1	35428.0	35428	35428	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	1305172	1	1305172.0	1305172	1305172	addVectorsInto(float*, float*, float*, int)

Number of blocks: 16, Threads per block: 256, Time Required: 1309816

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.4	253585855	3	84528618.3	3258	253574212	cudaMallocManaged
0.5	1312041	1	1312041.0	1312041	1312041	cudaDeviceSynchronize
0.1	250346	3	83448.7	10620	208049	cudaFree
0.0	40880	1	40880.0	40880	40880	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	1309816	1	1309816.0	1309816	1309816	addVectorsInto(float*, float*, float*, int)

Number of blocks: 32, Threads per block: 256, Time Required: 1406457

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.4	257983525	3	85994508.3	3841	257967279	cudaMallocManaged
0.5	1410940	1	1410940.0	1410940	1410940	cudaDeviceSynchronize
0.1	243629	3	81209.7	9975	202158	cudaFree
0.0	29226	1	29226.0	29226	29226	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	1406457	1	1406457.0	1406457	1406457	addVectorsInto(float*, float*, float*, int)

Number of blocks: 8, Threads per block: 512, Time Required: 1567085

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.3	258789829	3	86263276.3	6193	258766685	cudaMallocManaged
0.6	1587219	1	1587219.0	1587219	1587219	cudaDeviceSynchronize
0.1	325685	3	108561.7	13740	275675	cudaFree
0.0	42591	1	42591.0	42591	42591	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	1567085	1	1567085.0	1567085	1567085	addVectorsInto(float*, float*, float*, int)

Number of blocks: 16, Threads per block: 512, Time Required: 1391095

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.4	302167922	3	100722640.7	7493	302136720	cudaMallocManaged
0.4	1351968	1	1351968.0	1351968	1351968	cudaDeviceSynchronize
0.2	470686	3	156895.3	20743	398779	cudaFree
0.0	135244	1	135244.0	135244	135244	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	1391095	1	1391095.0	1391095	1391095	addVectorsInto(float*, float*, float*, int)

Number of blocks: 32, Threads per block: 512, Time Required: 1384345

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.3	249471482	3	83157160.7	3935	249453873	cudaMallocManaged
0.6	1388012	1	1388012.0	1388012	1388012	cudaDeviceSynchronize
0.1	234438	3	78146.0	10113	191420	cudaFree
0.0	30378	1	30378.0	30378	30378	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	1384345	1	1384345.0	1384345	1384345	addVectorsInto(float*, float*, float*, int)

N= 2<<24**Serial Execution Time: 2117915980**

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
88.6	2175923366	1	2175923366.0	2175923366	2175923366	cudaDeviceSynchronize
10.1	249249035	3	83083011.7	18262	249160342	cudaMallocManaged
0.9	21165881	3	7055293.7	6105795	8534822	cudaFree
0.4	10211381	3	3403793.7	7661	10050536	cudaMemPrefetchAsync
0.0	38289	1	38289.0	38289	38289	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	2117915980	1	2117915980.0	2117915980	2117915980	addVectorsInto(float*, float*, float*, int)

Parallel Execution Time:

Number of Blocks	Threads per Block	Time (in ns)	Speedup
8	256	10659168	198.6943
16	256	5827432	363.4389
32	256	3360477	630.2426
8	512	6078953	348.4014
16	512	3496636	605.7010
132	512	2148469	985.7791

Threads per block: 256, Number of blocks: 8, Time Required: 10659168

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
71.5	253195240	3	84398413.3	17526	253123280	cudaMallocManaged
12.7	45030822	1	45030822.0	45030822	45030822	cudaDeviceSynchronize
9.4	33370308	3	11123436.0	165062	22466204	cudaMemPrefetchAsync
6.3	22349575	3	7449858.3	6700811	8705098	cudaFree
0.0	36668	1	36668.0	36668	36668	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	10659168	1	10659168.0	10659168	10659168	addVectorsInto(float*, float*, float*, int)

Threads per block: 256, Number of blocks: 16, Time Required: 5827432

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
71.6	239709717	3	79903239.0	35708	239568662	cudaMallocManaged
11.2	37333789	1	37333789.0	37333789	37333789	cudaDeviceSynchronize
10.8	36235467	3	12078489.0	135113	22478737	cudaMemPrefetchAsync
6.4	21345459	3	7115153.0	6286294	8561414	cudaFree
0.0	38040	1	38040.0	38040	38040	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	5827432	1	5827432.0	5827432	5827432	addVectorsInto(float*, float*, float*, int)

Threads per block: 256, Number of blocks: 32, Time Required: 3360477

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
73.9	262692662	3	87564220.7	19089	262609446	cudaMallocManaged
10.7	37900928	3	12633642.7	158259	22605477	cudaMemPrefetchAsync
9.4	33240782	1	33240782.0	33240782	33240782	cudaDeviceSynchronize
6.0	21479652	3	7159884.0	6268936	8659407	cudaFree
0.0	41117	1	41117.0	41117	41117	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	3360477	1	3360477.0	3360477	3360477	addVectorsInto(float*, float*, float*, int)

Threads per block: 512, Number of blocks: 8, Time Required: 6078953

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
79.2	366247859	3	122082619.7	24337	366138814	cudaMallocManaged
9.0	41369623	1	41369623.0	41369623	41369623	cudaDeviceSynchronize
7.0	32442406	3	10814135.3	149806	22495493	cudaMemPrefetchAsync
4.8	22047012	3	7349004.0	6603656	8798783	cudaFree
0.0	39697	1	39697.0	39697	39697	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	6078953	1	6078953.0	6078953	6078953	addVectorsInto(float*, float*, float*, int)

Threads per block: 512, Number of blocks: 16, Time Required: 3496636

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
72.7	251162089	3	83720696.3	35468	250999154	cudaMallocManaged
17.7	61022632	1	61022632.0	61022632	61022632	cudaDeviceSynchronize
6.7	23073097	3	7691032.3	6258226	8622708	cudaFree
2.9	10126663	3	3375554.3	7856	9972133	cudaMemPrefetchAsync
0.0	36285	1	36285.0	36285	36285	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	3496636	1	3496636.0	3496636	3496636	addVectorsInto(float*, float*, float*, int)

Threads per block: 512, Number of blocks: 32, Time Required: 2148469

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
73.7	258462091	3	86154030.3	19090	258371602	cudaMallocManaged
10.7	37459080	1	37459080.0	37459080	37459080	cudaDeviceSynchronize
9.3	32494474	3	10831491.3	147396	22518420	cudaMemPrefetchAsync
6.4	22437186	3	7479062.0	6554815	8612458	cudaFree
0.0	40541	1	40541.0	40541	40541	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	2148469	1	2148469.0	2148469	2148469	addVectorsInto(float*, float*, float*, int)

N= 2<<10**Serial Execution Time: 266554**

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	427305624	3	142435208.0	7724	427266355	cudaMallocManaged
0.1	456341	3	152113.7	20364	404617	cudaMemPrefetchAsync
0.1	267087	1	267087.0	267087	267087	cudaDeviceSynchronize
0.0	153975	3	51325.0	17898	99035	cudaFree
0.0	39164	1	39164.0	39164	39164	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	266554	1	266554.0	266554	266554	addVectorsInto(float*, float*, float*, int)

Parallel Execution Time:

Number of Blocks	Threads per Block	Time (in ns)	Speedup
8	256	5440	48.9988
16	256	5216	51.1031
32	256	5344	49.8791
8	512	5664	47.0610
16	512	5472	48.7123
132	512	5665	47.0527

Threads per block: 256, Number of blocks: 8, Time Required: 5440

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	288048459	3	96016153.0	5528	288016841	cudaMallocManaged
0.1	413841	3	137947.0	18716	370145	cudaMemPrefetchAsync
0.0	131855	3	43951.7	15804	88695	cudaFree
0.0	36206	1	36206.0	36206	36206	cudaLaunchKernel
0.0	7962	1	7962.0	7962	7962	cudaDeviceSynchronize

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	5440	1	5440.0	5440	5440	addVectorsInto(float*, float*, float*, int)

Threads per block: 256, Number of blocks: 16, Time Required: 5216

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.7	274279608	3	91426536.0	3638	274261154	cudaMallocManaged
0.1	374808	1	374808.0	374808	374808	cudaDeviceSynchronize
0.1	355617	3	118539.0	18989	239004	cudaMemPrefetchAsync
0.0	108074	3	36024.7	11639	73272	cudaFree
0.0	62709	1	62709.0	62709	62709	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	5216	1	5216.0	5216	5216	addVectorsInto(float*, float*, float*, int)

Threads per block: 256, Number of blocks: 32, Time Required: 5344

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	254890908	3	84963636.0	5617	254856180	cudaMallocManaged
0.2	433121	3	144373.7	20363	387395	cudaMemPrefetchAsync
0.1	141081	3	47027.0	15056	94988	cudaFree
0.0	40716	1	40716.0	40716	40716	cudaLaunchKernel
0.0	8131	1	8131.0	8131	8131	cudaDeviceSynchronize

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	5344	1	5344.0	5344	5344	addVectorsInto(float*, float*, float*, int)

Threads per block: 512, Number of blocks: 8, Time Required: 5664

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.6	256999308	3	85666436.0	4304	256965354	cudaMallocManaged
0.1	377884	1	377884.0	377884	377884	cudaDeviceSynchronize
0.1	373425	3	124475.0	19140	250564	cudaMemPrefetchAsync
0.0	114766	3	38255.3	11556	75299	cudaFree
0.0	65032	1	65032.0	65032	65032	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	5664	1	5664.0	5664	5664	addVectorsInto(float*, float*, float*, int)

Threads per block: 512, Number of blocks: 16, Time Required: 5472

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	5472	1	5472.0	5472	5472	addVectorsInto(float*, float*, float*, int)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
56.6	8640	3	2880.0	2880	2880	[CUDA Unified Memory memcpy HtoD]
43.4	6622	3	2207.3	1632	3135	[CUDA Unified Memory memcpy DtoH]

Threads per block: 512, Number of blocks: 32, Time Required: 5665

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.7	304684255	3	101561418.3	8513	304635437	cudaMallocManaged
0.2	566573	3	188857.7	26426	504410	cudaMemPrefetchAsync
0.1	172369	3	57456.3	21393	112972	cudaFree
0.0	56456	1	56456.0	56456	56456	cudaLaunchKernel
0.0	10349	1	10349.0	10349	10349	cudaDeviceSynchronize

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	5665	1	5665.0	5665	5665	addVectorsInto(float*, float*, float*, int)

Conclusion:

From above graphs we can conclude that

1. For vector-vector addition problem for larger values of N time for parallel execution decreases exponentially as we increase number of threads
2. For vector-vector addition problem for smaller values of N time first decreases and then increases exponentially as we increase number of threads
3. Parallel execution is preferred for higher values of N

4. Implement N-Body Simulator using CUDA C. State and justify the speedup using different size of threads and blocks.

```
#include <math.h>
#include <stdio.h>
#include <stdlib.h>
#include "timer.h"
#include "files.h"
#define SOFTENING 1e-9f

/*
 * Each body contains x, y, and z coordinate positions,
 * as well as velocities in the x, y, and z directions.
 */

typedef struct { float x, y, z, vx, vy, vz; } Body;

/*
 * Calculate the gravitational impact of all bodies in the system
 * on all others.
 */
```

```

__global__ void bodyForce(Body *p, float dt, int N) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;

    if (tid < N) {
        float Fx = 0, Fy = 0, Fz = 0;
        for (int i = 0; i < N; i++) {
            float dx = p[i].x - p[tid].x;
            float dy = p[i].y - p[tid].y;
            float dz = p[i].z - p[tid].z;
            float distSqr = dx*dx + dy*dy + dz*dz + SOFTENING;
            float invDist = rsqrtf(distSqr);
            float invDist3 = invDist * invDist * invDist;
            Fx += dx * invDist3;
            Fy += dy * invDist3;
            Fz += dz * invDist3;
        }
        p[tid].vx += dt*Fx;
        p[tid].vy += dt*Fy;
        p[tid].vz += dt*Fz;
    }
}

```

```

int main(const int argc, const char** argv) {

```

```

    // The assessment will test against both 2<11 and 2<15.

```

```

    // Feel free to pass the command line argument 15 when you generate ./nbody report files

```

```

    int nBodies = 2<<11;

```

```
if (argc > 1) nBodies = 2<<atoi(argv[1]);
```

```
// The assessment will pass hidden initialized values to check for correctness.
```

```
// You should not make changes to these files, or else the assessment will not work.
```

```
const char * initialized_values;
```

```
const char * solution_values;
```

```
if (nBodies == 2<<11) {
```

```
    initialized_values = "09-nbody/files/initialized_4096";
```

```
    solution_values = "09-nbody/files/solution_4096";
```

```
} else { // nBodies == 2<<15
```

```
    initialized_values = "09-nbody/files/initialized_65536";
```

```
    solution_values = "09-nbody/files/solution_65536";
```

```
}
```

```
if (argc > 2) initialized_values = argv[2];
```

```
if (argc > 3) solution_values = argv[3];
```

```
const float dt = 0.01f; // Time step
```

```
const int nIters = 10; // Simulation iterations
```

```
int bytes = nBodies * sizeof(Body);
```

```
float *buf;
```

```
//buf = (float *)malloc(bytes);
```

```
cudaMallocManaged(&buf, bytes);
```

```
Body *p = (Body*)buf;
```



```

read_values_from_file(initialized_values, buf, bytes);

double totalTime = 0.0;

/*
 * This simulation will run for 10 cycles of time, calculating gravitational
 * interaction amongst bodies, and adjusting their positions to reflect.
 */

for (int iter = 0; iter < nIters; iter++) {
    StartTimer();
/*
 * You will likely wish to refactor the work being done in `bodyForce`,
 * and potentially the work to integrate the positions.
 */

    //int threads_per_block = 128;
    //int number_of_blocks = (nBodies / threads_per_block);
    int threads_per_block = 1;
    int number_of_blocks = 1;

    bodyForce <<< number_of_blocks, threads_per_block >>> ( p, dt, nBodies );
    //bodyForce<<<>>>(p, dt, nBodies); // compute interbody forces
    cudaDeviceSynchronize();
/*
 * This position integration cannot occur until this round of `bodyForce` has completed.
 * Also, the next round of `bodyForce` cannot begin until the integration is complete.
 */

```

```
for (int i = 0 ; i < nBodies; i++) { // integrate position
    p[i].x += p[i].vx*dt;
    p[i].y += p[i].vy*dt;
    p[i].z += p[i].vz*dt;
}

const double tElapsed = GetTimer() / 1000.0;
totalTime += tElapsed;
}

double avgTime = totalTime / (double)(nIters);
float billionsOfOpsPerSecond = 1e-9 * nBodies * nBodies / avgTime;
write_values_to_file(solution_values, buf, bytes);
// You will likely enjoy watching this value grow as you accelerate the application,
// but beware that a failure to correctly synchronize the device might result in
// unrealistically high values.
printf("%0.3f Billion Interactions / second\n", billionsOfOpsPerSecond);

cudaFree(buf);
}
```

N=2<<11**Serial Time Execution: 7951112**

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
96.8	244886108	1	244886108.0	244886108	244886108	cudaMallocManaged
3.2	7982348	10	798234.8	726906	1121157	cudaDeviceSynchronize
0.0	109377	1	109377.0	109377	109377	cudaFree
0.0	102904	10	10290.4	6602	33009	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	7951112	10	795111.2	723736	1118516	bodyForce(Body*, float, int)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
53.2	255563	46	5555.7	2175	11616	[CUDA Unified Memory memcpy HtoD]
46.8	225169	60	3752.8	1439	15200	[CUDA Unified Memory memcpy DtoH]

Parallel Execution Time:

Number of Blocks	Threads per Block	Time (in ns)	Speedup
2	128	8899145	0.8934
2	1024	20392480	0.3899
32	128	8958927	0.8875
32	1024	18692193	0.4253

Number of blocks: 2, Threads per block: 128, Execution Time: 8899145

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
96.5	250657886	1	250657886.0	250657886	250657886	cudaMallocManaged
3.4	8932173	10	893217.3	812619	1219339	cudaDeviceSynchronize
0.0	111204	10	11120.4	6712	33727	cudaLaunchKernel
0.0	110976	1	110976.0	110976	110976	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	8899145	10	889914.5	808952	1216756	bodyForce(Body*, float, int)

Number of blocks: 2, Threads per block: 1024, Execution Time: 20392480

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
92.6	260060741	1	260060741.0	260060741	260060741	cudaMallocManaged
7.3	20438083	10	2043808.3	1783428	3936978	cudaDeviceSynchronize
0.1	191800	10	19180.0	10666	34932	cudaLaunchKernel
0.0	113511	1	113511.0	113511	113511	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	20392480	10	2039248.0	1780751	3927867	bodyForce(Body*, float, int)

Number of blocks: 32, Threads per block: 128, Execution Time: 8958927

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
96.8	286120581	1	286120581.0	286120581	286120581	cudaMallocManaged
3.0	8876571	10	887657.1	814539	1069596	cudaDeviceSynchronize
0.1	351827	10	35182.7	12872	127660	cudaLaunchKernel
0.1	258560	1	258560.0	258560	258560	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	8958927	10	895892.7	849304	1138966	bodyForce(Body*, float, int)

Number of blocks: 32, Threads per block: 1024, Execution Time: 18692193

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
92.7	241462644	1	241462644.0	241462644	241462644	cudaMallocManaged
7.2	18727069	10	1872706.9	1790195	2501846	cudaDeviceSynchronize
0.1	133325	10	13332.5	8815	35724	cudaLaunchKernel
0.0	123531	1	123531.0	123531	123531	cudaFree
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	18692193	10	1869219.3	1785937	2498186	bodyForce(Body*, float, int)

N=2<<15

Serial Execution Time: 109233415

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
69.9	255879749	1	255879749.0	255879749	255879749	cudaMallocManaged
29.9	109301730	10	10930173.0	10846567	11419808	cudaDeviceSynchronize
0.1	467823	1	467823.0	467823	467823	cudaFree
0.1	291833	10	29183.3	22028	43391	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	109233415	10	10923341.5	10841446	11413722	bodyForce(Body*, float, int)

Parallel Execution Time:

Number of Blocks	Threads per Block	Time (in ns)	Speedup
2	128	120612160	0.9056
2	1024	229499125	0.0888
32	128	116887945	0.9345
32	1024	264315658	0.4132

Number of blocks: 2, Threads per block: 128, Execution Time: 120612160

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
68.7	265590539	1	265590539.0	265590539	265590539	cudaMallocManaged
31.2	120673622	10	12067362.2	11771548	12628154	cudaDeviceSynchronize
0.1	329526	1	329526.0	329526	329526	cudaFree
0.1	260968	10	26096.8	22295	36503	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	120612160	10	12061216.0	11764844	12621883	bodyForce(Body*, float, int)

Number of blocks: 2, Threads per block: 1024, Execution Time: 229499125

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
53.3	262847173	1	262847173.0	262847173	262847173	cudaMallocManaged
46.5	229594645	10	22959464.5	19688324	27692528	cudaDeviceSynchronize
0.1	437803	10	43780.3	38819	51465	cudaLaunchKernel
0.1	342831	1	342831.0	342831	342831	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	229499125	10	22949912.5	19680193	27680088	bodyForce(Body*, float, int)

Number of blocks: 32, Threads per block: 128, Execution Time: 116887945

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
78.2	423221770	1	423221770.0	423221770	423221770	cudaMallocManaged
21.6	116942222	10	11694222.2	11591034	12261782	cudaDeviceSynchronize
0.1	401549	10	40154.9	21260	98015	cudaLaunchKernel
0.1	311617	1	311617.0	311617	311617	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	116887945	10	11688794.5	11585674	12261083	bodyForce(Body*, float, int)

Number of blocks: 32, Threads per block: 1024, Execution Time: 264315658

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
50.7	263706704	10	26370670.4	25506033	27620883	cudaDeviceSynchronize
49.1	255184018	1	255184018.0	255184018	255184018	cudaMallocManaged
0.1	511319	10	51131.9	34002	107947	cudaLaunchKernel
0.1	361911	1	361911.0	361911	361911	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	264315658	10	26431565.8	25953208	27610805	bodyForce(Body*, float, int)

Conclusion:

1. From above graph we can conclude that
2. For nbody problem parallel execution time first decreases exponentially and then increases as we increase number of threads