

Final Year B. Tech., Sem VII 2022-23

High Performance Computing Lab

PRN: 2020BTECS00206

Full Name: SAYALI YOGESH DESAI

Batch: B4

Assignment No. 10

1. Implement Matrix-matrix Multiplication using global memory in CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

```
#include <stdio.h>
```

```
void initWith(float num, float *a, int SIZE)
```

```
{
```

```
    for(int i = 0; i < SIZE; ++i)
```

```
    {
```

```
        a[i] = num;
```

```
    }
```

```
}
```

```
__global__
```

```
void matrixMultiply(float *result, float *a, float *b, int N, int SIZE)
```

```
{
```

```
    int start = blockIdx.x * blockDim.x + threadIdx.x;
```

```
    int stride = gridDim.x * blockDim.x;
```

```
    for(int i = start; i < SIZE; i += stride)
```

```
{
    int row = i / N;
    float sum = 0
    for (int j = 0; j < N; j++)
    {
        sum += a[row * N + j] * b[N * j + row];
    }
    result[i] = sum;
}

void checkElementsAre(float target, float *array, int SIZE)
{
    for(int i = 0; i < SIZE; i++)
    {
        if(array[i] != target)
        {
            printf("FAIL: array[%d] - %0.0f does not equal %0.0f\n", i, array[i], target);
            exit(1);
        }
    }
    printf("SUCCESS! All values multiplied correctly.\n");
}

int main()
{
    const int N = 1024;
    const int SIZE = N * N; // square matrix
    size_t size = SIZE * sizeof(float);
```

```
float *a;
```

```
float *b;
```

```
float *c;
```

```
cudaMallocManaged(&a, size);
```

```
cudaMallocManaged(&b, size);
```

```
cudaMallocManaged(&c, size);
```

```
initWith(3, a, SIZE);
```

```
initWith(4, b, SIZE);
```

```
initWith(0, c, SIZE);
```

```
matrixMultiply<<<1, 1>>>(c, a, b, N, SIZE);
```

```
cudaDeviceSynchronize();
```

```
checkElementsAre(12288, c, SIZE);
```

```
cudaFree(a);
```

```
cudaFree(b);
```

```
cudaFree(c);
```

```
}
```

8*8 Matrix:**Serial Execution Time: 399933ns**

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	333656825	3	111218941.7	5627	333622273	cudaMallocManaged
0.1	402410	1	402410.0	402410	402410	cudaDeviceSynchronize
0.0	115065	3	38355.0	11233	71366	cudaFree
0.0	38982	1	38982.0	38982	38982	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	399933	1	399933.0	399933	399933	matrixMultiply(float*, float*, float*, int, int)

Parallel Execution Time:

Number of blocks	Thread per blocks	Time (in ns)	Speedup
16	512	268958	1.4869
16	1024	334204	1.1966
32	512	374204	1.0687
32	1024	332125	1.2041

Number of blocks: 16, Thread per blocks: 512, Execution Time: 268958

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	430775041	3	143591680.3	12051	430744161	cudaMallocManaged
0.1	310387	3	103462.3	22697	233921	cudaFree
0.1	267937	1	267937.0	267937	267937	cudaDeviceSynchronize
0.0	53001	1	53001.0	53001	53001	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	268958	1	268958.0	268958	268958	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 16, Thread per blocks: 1024, Execution Time: 334204

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	407981405	3	135993801.7	6173	407950302	cudaMallocManaged
0.1	337501	1	337501.0	337501	337501	cudaDeviceSynchronize
0.0	114739	3	38246.3	11446	72049	cudaFree
0.0	30932	1	30932.0	30932	30932	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	334204	1	334204.0	334204	334204	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 32, Thread per blocks: 512, Execution Time: 374204

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	242384591	3	80794863.7	5882	242348869	cudaMallocManaged
0.2	377506	1	377506.0	377506	377506	cudaDeviceSynchronize
0.0	114198	3	38066.0	11038	68457	cudaFree
0.0	36897	1	36897.0	36897	36897	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	374204	1	374204.0	374204	374204	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 32, Thread per blocks: 1024, Execution Time: 332125

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	430959046	3	143653015.3	11084	430898885	cudaMallocManaged
0.1	335709	1	335709.0	335709	335709	cudaDeviceSynchronize
0.1	334428	3	111476.0	23334	252894	cudaFree
0.0	65830	1	65830.0	65830	65830	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	332125	1	332125.0	332125	332125	matrixMultiply(float*, float*, float*, int, int)

1024*1024 Matrix:**Serial Execution Time: 17572204446ns**

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
98.6	17572214481	1	17572214481.0	17572214481	17572214481	cudaDeviceSynchronize
1.4	249008507	3	83002835.7	11156	248965441	cudaMallocManaged
0.0	1012534	3	337511.3	252116	429267	cudaFree
0.0	44857	1	44857.0	44857	44857	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	17572204446	1	17572204446.0	17572204446	17572204446	matrixMultiply(float*, float*, float*, int, int)

Parallel Execution Time:

Number of blocks	Thread per blocks	Time (in ns)	Speedup
16	512	22561818	778.84
16	1024	21109170	832.44
32	512	15056306	1167.09
32	1024	14094717	1246.72

Number of blocks: 16, Thread per blocks: 512, Execution Time: 22561818ns

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
94.7	428180509	3	142726836.3	19497	428090610	cudaMallocManaged
5.0	22570901	1	22570901.0	22570901	22570901	cudaDeviceSynchronize
0.3	1395024	3	465008.0	392268	511781	cudaFree
0.0	58460	1	58460.0	58460	58460	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	22561818	1	22561818.0	22561818	22561818	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 16, Thread per blocks: 1024, Execution Time: 21109170ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
94.2	364405996	3	121468665.3	12181	364348764	cudaMallocManaged
5.5	21119573	1	21119573.0	21119573	21119573	cudaDeviceSynchronize
0.4	1410478	3	470159.3	352370	556149	cudaFree
0.0	42436	1	42436.0	42436	42436	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	21109170	1	21109170.0	21109170	21109170	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 32, Thread per blocks: 512, Execution Time: 15056306ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
94.1	266747002	3	88915667.3	16877	266668877	cudaMallocManaged
5.3	15065454	1	15065454.0	15065454	15065454	cudaDeviceSynchronize
0.5	1469173	3	489724.3	386699	602450	cudaFree
0.0	49207	1	49207.0	49207	49207	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	15056306	1	15056306.0	15056306	15056306	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 32, Thread per blocks: 1024, Execution Time: 14094717ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
96.0	389416495	3	129805498.3	24163	389296767	cudaMallocManaged
3.5	14002462	1	14002462.0	14002462	14002462	cudaDeviceSynchronize
0.5	2126085	3	708695.0	532197	926545	cudaFree
0.0	185419	1	185419.0	185419	185419	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	14094717	1	14094717.0	14094717	14094717	matrixMultiply(float*, float*, float*, int, int)

2. Implement Matrix-Matrix Multiplication using shared memory in CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

```
#include <stdio.h>

void initWith(float num, float *a, int SIZE)

{

    for(int i = 0; i < SIZE; ++i)

    {

        a[i] = num;

    }

}

__global__

void matrixMultiply(float *result, float *a, float *b, int N, int SIZE)

{

    __shared__ int stride;

    if (threadIdx.x == 0)

        stride = gridDim.x * blockDim.x;

    __syncthreads();

    int start = blockIdx.x * blockDim.x + threadIdx.x;

    for(int i = start; i < SIZE; i += stride)

    {

        int row = i / N;
```



```
float sum = 0;

for (int j = 0; j < N; j++)

{

    sum += a[row * N + j] * b[N * j + row];

}

result[i] = sum;

}

}

void checkElementsAre(float target, float *array, int SIZE)

{

for(int i = 0; i < SIZE; i++)

{

    if(array[i] != target)

    {

        printf("FAIL: array[%d] - %0.0f does not equal %0.0f\n", i, array[i], target);

        exit(1);

    }

}

printf("SUCCESS! All values multiplied correctly.\n");

}

int main()

{
```

```
const int N = 1024;

const int SIZE = N * N; // square matrix

size_t size = SIZE * sizeof(float);


float *a;

float *b;

float *c;


cudaMallocManaged(&a, size);

cudaMallocManaged(&b, size);

cudaMallocManaged(&c, size);


initWith(3, a, SIZE);

initWith(4, b, SIZE);

initWith(0, c, SIZE);

matrixMultiply<<<1, 1>>>(c, a, b, N, SIZE);

cudaDeviceSynchronize();

checkElementsAre(12288, c, SIZE);

cudaFree(a);

cudaFree(b);

cudaFree(c);

}
```

8*8 Matrix:**Serial Execution Time: 549241ns**

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.7	265548138	3	88516046.0	8292	265502071	cudaMallocManaged
0.2	550937	1	550937.0	550937	550937	cudaDeviceSynchronize
0.1	152321	3	50773.7	16226	96754	cudaFree
0.0	53871	1	53871.0	53871	53871	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	549241	1	549241.0	549241	549241	matrixMultiply(float*, float*, float*, int, int)

Parallel Execution Time:

Number of blocks	Thread per blocks	Time (in ns)	Speedup
16	512	400923	1.3699
16	1024	380507	1.4434
32	512	481114	1.1416
32	1024	388795	1.4126

Number of blocks: 16, Thread per blocks: 512, Execution Time: 400923ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	358975781	3	119658593.7	9435	358928286	cudaMallocManaged
0.1	401915	1	401915.0	401915	401915	cudaDeviceSynchronize
0.0	148646	3	49548.7	15696	90886	cudaFree
0.0	58320	1	58320.0	58320	58320	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	400923	1	400923.0	400923	400923	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 16, Thread per blocks: 1024, Execution Time: 380507ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	244953469	3	81651156.3	6348	244912256	cudaMallocManaged
0.2	382983	1	382983.0	382983	382983	cudaDeviceSynchronize
0.1	125462	3	41820.7	11562	77136	cudaFree
0.0	43366	1	43366.0	43366	43366	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	380507	1	380507.0	380507	380507	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 32, Thread per blocks: 512, Execution Time: 481114ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.7	252737992	3	84245997.3	6972	252696136	cudaMallocManaged
0.2	482562	1	482562.0	482562	482562	cudaDeviceSynchronize
0.1	153475	3	51158.3	15797	94368	cudaFree
0.0	49459	1	49459.0	49459	49459	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	481114	1	481114.0	481114	481114	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 32, Thread per blocks: 1024, Execution Time: 388795ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	396923132	3	132307710.7	12306	396855209	cudaMallocManaged
0.1	388135	1	388135.0	388135	388135	cudaDeviceSynchronize
0.1	301673	3	100557.7	22093	221470	cudaFree
0.0	65971	1	65971.0	65971	65971	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	388795	1	388795.0	388795	388795	matrixMultiply(float*, float*, float*, int, int)

1024*1024 Matrix:**Serial Execution Time: 17346175067ns**

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
97.6	17346186443	1	17346186443.0	17346186443	17346186443	cudaDeviceSynchronize
2.4	420614935	3	140204978.3	23484	420540161	cudaMallocManaged
0.0	989852	3	329950.7	242389	435729	cudaFree
0.0	52964	1	52964.0	52964	52964	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	17346175067	1	17346175067.0	17346175067	17346175067	matrixMultiply(float*, float*, float*, int, in

Parallel Execution Time:

Number of blocks	Thread per blocks	Time (in ns)	Speedup
16	512	22288180	778.26
16	1024	20991075	826.35
32	512	15596802	1049.72
32	1024	14524471	1112.16

Number of blocks: 16, Thread per blocks: 512, Execution Time: 22288180ns

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
91.3	248057568	3	82685856.0	12426	247988053	cudaMallocManaged
8.2	22298286	1	22298286.0	22298286	22298286	cudaDeviceSynchronize
0.5	1333775	3	444591.7	363907	524273	cudaFree
0.0	50506	1	50506.0	50506	50506	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	22288180	1	22288180.0	22288180	22288180	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 16, Thread per blocks: 1024, Execution Time: 20991075ns

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
91.9	250078234	3	83359411.3	11550	250029665	cudaMallocManaged
7.7	21000813	1	21000813.0	21000813	21000813	cudaDeviceSynchronize
0.4	1079357	3	359785.7	245469	531308	cudaFree
0.0	58180	1	58180.0	58180	58180	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	20991075	1	20991075.0	20991075	20991075	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 32, Thread per blocks: 512, Execution Time: 15596802ns

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
93.5	244222577	3	81407525.7	17141	244136369	cudaMallocManaged
6.0	15605230	1	15605230.0	15605230	15605230	cudaDeviceSynchronize
0.5	1349153	3	449717.7	361335	547870	cudaFree
0.0	58678	1	58678.0	58678	58678	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	15596802	1	15596802.0	15596802	15596802	matrixMultiply(float*, float*, float*, int, int)

Number of blocks: 32, Thread per blocks: 1024, Execution Time: 16524471ns

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
93.6	256217265	3	85405755.0	12154	256163814	cudaMallocManaged
6.0	16532739	1	16532739.0	16532739	16532739	cudaDeviceSynchronize
0.3	942942	3	314314.0	240919	395541	cudaFree
0.0	50569	1	50569.0	50569	50569	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	16524471	1	16524471.0	16524471	16524471	matrixMultiply(float*, float*, float*, int, int)

3. Implement Prefix sum using CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

```
#include <stdio.h>

void initWith(float val, float *arr, int N)

{

    for (int i = 0; i < N; i++)

    {

        arr[i] = val;

    }

}

__global__

void prefixSum(float *arr, float *res, float *ptemp, float* ttemp, int N)

{

    int threadId = blockIdx.x * blockDim.x + threadIdx.x;

    int totalThreads = gridDim.x * blockDim.x;

    int elementsPerThread = ceil(1.0 * N / totalThreads);

    int start = threadId * elementsPerThread;

    int count = 0;

    float *sums = new float[elementsPerThread];

    float sum = 0;

    for (int i = start; i < N && count < elementsPerThread; i++, count++) {
```

```
    sum += arr[i];

    sums[count] = sum;

}

float localSum;

if (count)

    localSum = sums[count - 1];

else

    localSum = 0;

ptemp[threadId] = localSum;

ttemp[threadId] = localSum;

__syncthreads();

if (totalThreads == 1) {

    for (int i = 0; i < N; i++)

        res[i] = sums[i];

} else {

    int d = 0; // log2(totalThreads)

    int x = totalThreads;

    while (x > 1) {

        d++;

        x = x >> 1;

    }
```



```

x = 1;

for (int i = 0; i < 2*d; i++) {

    int tsum = ttemp[threadId];

    __syncthreads();

    int newId = threadId / x;

    if (newId % 2 == 0) {

        int nextId = threadId + x;

        ptemp[nextId] += tsum;

        ttemp[nextId] += tsum;

    } else {

        int nextId = threadId - x;

        ttemp[nextId] += tsum;

    }

    x = x << 1;

}

__syncthreads();

float diff = ptemp[threadId] - localSum;

for (int i = start, j = 0; i < N && j < count; i++, j++) {

    res[i] = sums[j] + diff;

}

}

}

```

```
void checkRes(float *arr, float *res, int N, float *ptemp, float* ttemp)

{

    float sum = 0;

    for (int i = 0; i < N; i++)

    {

        sum += arr[i];

        if (sum != res[i])

        {

            printf("FAIL: res[%d] - %0.0f does not equal %0.0f\n", i, res[i], sum);

            exit(1);

        }

    }

    printf("SUCCESS! All prefix sums added correctly.\n");

}

int main()

{

    const int N = 1000000;

    size_t size = N * sizeof(float);

    float *arr;

    float *res;

    cudaMallocManaged(&arr, size);

    cudaMallocManaged(&res, size);
```

```
initWith(2, arr, N);

initWith(0, res, N);

int blocks = 1;

int threadsPerBlock = 1;

int totalThreads = blocks * threadsPerBlock;

float *ptemp;

float *ttemp;

cudaMallocManaged(&ptemp, totalThreads * sizeof(float));

cudaMallocManaged(&ttemp, totalThreads * sizeof(float));


prefixSum<<<blocks, threadsPerBlock>>>(arr, res, ptemp, ttemp, N);

cudaDeviceSynchronize();

checkRes(arr, res, N, ptemp, ttemp);


cudaFree(arr);

cudaFree(res);

cudaFree(ttemp);

cudaFree(ptemp);

}
```

Element 10^6**Serial Execution Time: 208767146ns**

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
55.0	257016448	4	64254112.0	6366	256904067	cudaMallocManaged
44.7	208779362	1	208779362.0	208779362	208779362	cudaDeviceSynchronize
0.2	839288	4	209822.0	18265	426965	cudaFree
0.1	338030	1	338030.0	338030	338030	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	208767146	1	208767146.0	208767146	208767146	prefixSum(float*, float*, float*, float*, int)

Parallel Execution Time:

Number of Blocks	Threads per Block	Time(ns)	Speedup
1	32	24535074	8.5089
1	64	18895276	11.0486
1	128	17276600	12.0838

Number of Blocks: 1, Thread per Blocks: 32, Execution Time: 24535074ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
91.2	268983360	4	67245840.0	8158	268851324	cudaMallocManaged
8.3	24545786	1	24545786.0	24545786	24545786	cudaDeviceSynchronize
0.3	846762	4	211690.5	17543	432428	cudaFree
0.2	457811	1	457811.0	457811	457811	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	24535074	1	24535074.0	24535074	24535074	prefixSum(float*, float*, float*, float*, int)

Number of Blocks: 1, Thread per Blocks: 64, Execution Time: 18895276ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
95.4	426974920	4	106743730.0	13060	426801429	cudaMallocManaged
4.2	18897120	1	18897120.0	18897120	18897120	cudaDeviceSynchronize
0.2	863297	4	215824.3	17095	453620	cudaFree
0.1	649967	1	649967.0	649967	649967	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	18895276	1	18895276.0	18895276	18895276	prefixSum(float*, float*, float*, float*, int)

Number of Blocks: 1, Thread per Blocks: 128, Execution Time: 17276600ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
93.4	273746331	4	68436582.8	4019	273655813	cudaMallocManaged
5.9	17285891	1	17285891.0	17285891	17285891	cudaDeviceSynchronize
0.6	1611511	4	402877.8	29374	803088	cudaFree
0.1	305410	1	305410.0	305410	305410	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	17276600	1	17276600.0	17276600	17276600	prefixSum(float*, float*, float*, float*, int)

Parallel Execution Time:

Elements	Number of Blocks	Threads per Block	Time(ns)
10 ³	1	64	548220
10 ⁴	1	64	1462645
10 ⁵	1	64	2765066
10 ⁶	1	64	19337795

10^3, Execution time: 548220ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.6	259781846	4	64945461.5	4548	259750212	cudaMallocManaged
0.2	549957	1	549957.0	549957	549957	cudaDeviceSynchronize
0.1	308813	1	308813.0	308813	308813	cudaLaunchKernel
0.0	114302	4	28575.5	10799	59995	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	548220	1	548220.0	548220	548220	prefixSum(float*, float*, float*, float*, int)

10^4, Execution time: 1462645ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.4	330582419	4	82645604.8	6766	330549315	cudaMallocManaged
0.4	1464154	1	1464154.0	1464154	1464154	cudaDeviceSynchronize
0.1	385842	1	385842.0	385842	385842	cudaLaunchKernel
0.0	156093	4	39023.3	15535	89045	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	1462645	1	1462645.0	1462645	1462645	prefixSum(float*, float*, float*, float*, int)

10^5, Execution time: 2765066ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
98.8	292890571	4	73222642.8	7807	292835338	cudaMallocManaged
0.9	2769867	1	2769867.0	2769867	2769867	cudaDeviceSynchronize
0.1	419075	1	419075.0	419075	419075	cudaLaunchKernel
0.1	287849	4	71962.3	14712	208806	cudaFree

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	2765066	1	2765066.0	2765066	2765066	prefixSum(float*, float*, float*, float*, int)

10⁶, Execution time: 19337795ns

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
92.5	254590340	4	63647585.0	5837	254485798	cudaMallocManaged
7.0	19344471	1	19344471.0	19344471	19344471	cudaDeviceSynchronize
0.3	864252	4	216063.0	17372	448916	cudaFree
0.1	331967	1	331967.0	331967	331967	cudaLaunchKernel

CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	19337795	1	19337795.0	19337795	19337795	prefixSum(float*, float*, float*, float*, int)

4. Implement 2D Convolution using shared memory using CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

```
#include <stdio.h>

#define MASK_DIM 7

#define MASK_OFFSET (MASK_DIM / 2)

__constant__ int mask[7 * 7];

__global__ void convolution_2d(int *matrix, int *result, int N)
{
    // Calculate the global thread positions

    int row = blockIdx.y * blockDim.y + threadIdx.y;

    int col = blockIdx.x * blockDim.x + threadIdx.x;

    // Starting index for calculation

    int start_r = row - MASK_OFFSET;

    int start_c = col - MASK_OFFSET;
```

```
// Temp value for accumulating the result

int temp = 0;

// Iterate over all the rows

for (int i = 0; i < MASK_DIM; i++)

{

    // Go over each column

    for (int j = 0; j < MASK_DIM; j++)

    {

        // Range check for rows

        if ((start_r + i) >= 0 && (start_r + i) < N)

        {

            // Range check for columns

            if ((start_c + j) >= 0 && (start_c + j) < N)

            {

                // Accumulate result

                temp += matrix[(start_r + i) * N + (start_c + j)] * mask[i * MASK_DIM + j];

            }

        }

    }

}

// Write back the result

result[row * N + col] = temp;
```



```
}

void init_matrix(int *m, int n)

{
    for (int i = 0; i < n; i++)

    {
        for (int j = 0; j < n; j++)

        {
            m[n * i + j] = rand() % 100;

        }

    }
}

void verify_result(int *m, int *mask, int *result, int N)

{
    int temp;

    int offset_r;

    int offset_c;

    // Go over each row

    for (int i = 0; i < N; i++)

    {

        // Go over each column

        for (int j = 0; j < N; j++)

        {
```

```
// Reset the temp variable

temp = 0;

// Go over each mask row

for (int k = 0; k < MASK_DIM; k++)

{

    // Update offset value for row

    offset_r = i - MASK_OFFSET + k;

    // Go over each mask column

    for (int l = 0; l < MASK_DIM; l++)

    {

        // Update offset value for column

        offset_c = j - MASK_OFFSET + l;

        // Range checks if we are hanging off the matrix

        if (offset_r >= 0 && offset_r < N)

        {

            if (offset_c >= 0 && offset_c < N)

            {

                // Accumulate partial results

                temp += m[offset_r * N + offset_c] * mask[k * MASK_DIM + l];

            }

        }

    }

}
```

```
    }

    // Fail if the results don't match

    if (result[i * N + j] != temp)

    {

        printf("Check failed");

        return;

    }

}

}

}

}

int main()

{

    int N = 1 << 10; // 2^10

    size_t bytes_n = N * N * sizeof(int);

    size_t bytes_m = MASK_DIM * MASK_DIM * sizeof(int);

    int *matrix;

    int *result;

    int *h_mask;
```

```
cudaMallocManaged(&matrix, bytes_n);

cudaMallocManaged(&result, bytes_n);

cudaMallocManaged(&h_mask, bytes_m);


init_matrix(matrix, N);

init_matrix(mask, MASK_DIM);


cudaMemcpyToSymbol(mask, h_mask, bytes_m);


// Calculate grid dimensions

int THREADS = 1;

int BLOCKS = (N + THREADS - 1) / THREADS;


// Dimension launch arguments

dim3 block_dim(THREADS, THREADS);

dim3 grid_dim(BLOCKS, BLOCKS);


convolution_2d<<<grid_dim, block_dim>>>(matrix, result, N);


verify_result(matrix, h_mask, result, N);


printf("COMPLETED SUCCESSFULLY!");
```

```

cudaFree(matrix);

cudaFree(result);

cudaFree(h_mask);

return 0;

}

```

Execution Time:

Threads	$2^4 \times 2^4$	Speedup
Serial Execution:	506396	-
4	638810	0.7927
8	382717	1.3231
16	531931	0.9519
32	517883	0.9778

Threads	$2^5 \times 2^5$	Speedup
Serial Execution:	769561	-
4	564605	1.3630
8	452604	1.7002
16	439085	1.7526
32	400955	1.9193

Threads	$2^{10} \times 2^{10}$	Speedup
Serial Execution:	14161311	-
4	4514933	3.1365
8	4055855	3.4915
16	460935	30.7230
32	432174	32.7676

2^4

Thread=1, Execution Time: 506396ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	308293138	3	102764379.3	6116	308258978	cudaMallocManaged
0.2	536804	3	178934.7	13950	437958	cudaFree
0.0	47066	1	47066.0	47066	47066	cudaMemcpyToSymbol
0.0	32138	1	32138.0	32138	32138	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	506396	1	506396.0	506396	506396	convolution_2d(int*, int*, int)

Thread=4, Execution Time: 638810ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	322418643	3	107472881.0	7011	322380423	cudaMallocManaged
0.0	147496	3	49165.3	20221	89048	cudaFree
0.0	49048	1	49048.0	49048	49048	cudaMemcpyToSymbol
0.0	39713	1	39713.0	39713	39713	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	638810	1	638810.0	638810	638810	convolution_2d(int*, int*, int)

Thread=8, Execution Time: 382717ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	242726932	3	80908977.3	4722	242701907	cudaMallocManaged
0.2	371918	3	123972.7	10332	305877	cudaFree
0.0	37507	1	37507.0	37507	37507	cudaMemcpyToSymbol
0.0	29886	1	29886.0	29886	29886	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	382717	1	382717.0	382717	382717	convolution_2d(int*, int*, int)

Thread=16, Execution Time: 531931ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	284286035	3	94762011.7	5998	284252122	cudaMallocManaged
0.2	545902	3	181967.3	13703	445783	cudaFree
0.0	47575	1	47575.0	47575	47575	cudaMemcpyToSymbol
0.0	36702	1	36702.0	36702	36702	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	531931	1	531931.0	531931	531931	convolution_2d(int*, int*, int)

Thread=32, Execution Time: 517883ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.7	246534024	3	82178008.0	4375	246514057	cudaMallocManaged
0.2	558570	3	186190.0	10428	479990	cudaFree
0.0	35563	1	35563.0	35563	35563	cudaMemcpyToSymbol
0.0	24677	1	24677.0	24677	24677	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	517883	1	517883.0	517883	517883	convolution_2d(int*, int*, int)

2^5

Thread=1, Execution Time: 769561ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	431403704	3	143801234.7	7632	431363340	cudaMallocManaged
0.1	219817	3	73272.3	20505	121298	cudaFree
0.0	54509	1	54509.0	54509	54509	cudaMemcpyToSymbol
0.0	36541	1	36541.0	36541	36541	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	769561	1	769561.0	769561	769561	convolution_2d(int*, int*, int)

Thread=4, Execution Time: 404605ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	235963670	3	78654556.7	11132	235924019	cudaMallocManaged
0.1	278212	3	92737.3	19346	210995	cudaFree
0.1	123241	1	123241.0	123241	123241	cudaLaunchKernel
0.0	60773	1	60773.0	60773	60773	cudaMemcpyToSymbol

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	404605	1	404605.0	404605	404605	convolution_2d(int*, int*, int)

Thread=8, Execution Time: 452604ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	251829350	3	83943116.7	6699	251796602	cudaMallocManaged
0.0	124725	3	41575.0	10483	74268	cudaFree
0.0	44367	1	44367.0	44367	44367	cudaMemcpyToSymbol
0.0	35315	1	35315.0	35315	35315	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	452604	1	452604.0	452604	452604	convolution_2d(int*, int*, int)

Thread=16, Execution Time: 409085ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	272423022	3	90807674.0	6203	272385760	cudaMallocManaged
0.0	116661	3	38887.0	10985	68605	cudaFree
0.0	46890	1	46890.0	46890	46890	cudaMemcpyToSymbol
0.0	32157	1	32157.0	32157	32157	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	409085	1	409085.0	409085	409085	convolution_2d(int*, int*, int)

Thread=32, Execution Time: 560955ns

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	292932868	3	97644289.3	6992	292894524	cudaMallocManaged
0.1	151886	3	50628.7	14292	91421	cudaFree
0.0	49895	1	49895.0	49895	49895	cudaMemcpyToSymbol
0.0	38287	1	38287.0	38287	38287	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	560955	1	560955.0	560955	560955	convolution_2d(int*, int*, int)